

Bellabeat Case Study - Fitbit Data Analysis

Tim Foltz

2025-08-11

Contents

Executive Summary	1
Ask	2
Prepare	2
Process	3
Analyze	3
KPI Snapshot	3
Average Steps by Weekday	4
Steps vs Calories (Outliers Removed)	5
Steps \times Sleep Hours \rightarrow Mean Calories (Heatmap)	5
Steps \times Sleep Hours by Weekend/Weekday (Heatmap)	7
Active Hours by Weekday and Intensity	7
Weekend vs Weekday (Goal & Distribution)	8
Hourly Steps Heatmap	10
Share (Slide Deck Outline)	11
Act (Recommendations & Next Steps)	13
Limitations	14
Reproducibility	14

Executive Summary

Business Task. Analyze public Fitbit tracker data to find usage trends and translate them into marketing opportunities for Bellabeat products.

Key Findings. - Steps and calories are strongly related; most high-calorie days pair with moderate-high steps and ~6–9 hours of sleep. - Light activity dominates time budgets; very-active minutes are scarce. - Weekends and weekdays have similar odds of hitting 10k steps; weekends show wider spread in sleep hours. - Hourly patterns peak mid-day and early evening.

Recommendations. - Run a cross-week 10k Steps Challenge with streak badges and conversion tracking. - Trigger gentle-day routines when sleep < 7 hours (lighter goals, reminders). - Promote intensity-mix micro-goals (e.g., 2×10 minutes fairly active + 40 minutes lightly active).

Ask

- What trends exist in smart device usage?
 - How do those trends apply to Bellabeat customers?
 - How can they influence Bellabeat's marketing strategy?
-

Prepare

Data. Fitbit Fitness Tracker Dataset (Mobius/Kaggle, CC0).

Periods. 3.12.16–4.11.16 and 4.12.16–5.12.16.

Notes. Small convenience sample; missingness and non-wear present.

```
library(tidyverse)
library(lubridate)
library(fs)
library(janitor)
library(ggplot2)
library(scales)
library(viridis)
```

```
folder1 <- "data/mturkfitbit_export_3.12.16-4.11.16/Fitabase_Data_3.12.16-4.11.16"
folder2 <- "data/mturkfitbit_export_4.12.16-5.12.16/Fitabase_Data_4.12.16-5.12.16"
```

```
file_paths1 <- fs::dir_ls(folder1, glob = "*.csv")
file_paths2 <- fs::dir_ls(folder2, glob = "*.csv")
```

```
file_contents_1 <- setNames(lapply(file_paths1, readr::read_csv, show_col_types = FALSE), basename(file_paths1))
file_contents_2 <- setNames(lapply(file_paths2, readr::read_csv, show_col_types = FALSE), basename(file_paths2))
```

```
all_files <- union(names(file_contents_1), names(file_contents_2))
data <- lapply(all_files, function(fname) {
  file1 <- file_contents_1[[fname]]
  file2 <- file_contents_2[[fname]]
  if (!is.null(file1) && !is.null(file2)) bind_rows(file1, file2) else if (!is.null(file1)) file1 else file2
})
names(data) <- all_files
```

Process

```
daily_activity <- data[["dailyActivity_merged.csv"]] %>%
  janitor::clean_names() %>%
  mutate(
    activity_date = as.Date(activity_date, format = "%m/%d/%Y"),
    weekday = wday(activity_date, label = TRUE, abbr = FALSE, week_start = 1),
    total_active_minutes = very_active_minutes + fairly_active_minutes + lightly_active_minutes
  ) %>%
  distinct()

sleep_data <- data[["sleepDay_merged.csv"]] %>%
  janitor::clean_names() %>%
  mutate(sleep_day = as.Date(sleep_day, format = "%m/%d/%Y")) %>%
  distinct(id, sleep_day, .keep_all = TRUE)

activity_sleep <- daily_activity %>%
  left_join(sleep_data, by = c("id", "activity_date" = "sleep_day"))

daily_activity <- daily_activity %>%
  filter(!(total_steps == 0 & calories > 2500))
activity_sleep <- activity_sleep %>%
  filter(!(total_steps == 0 & calories > 2500))
```

Analyze

KPI Snapshot

```
kpis <- activity_sleep %>%
  mutate(
    total_hours_asleep = total_minutes_asleep / 60,
    total_active_hours = total_active_minutes / 60
  ) %>%
  summarise(
    avg_steps = mean(total_steps, na.rm = TRUE),
    median_steps = median(total_steps, na.rm = TRUE),
    avg_active_hours = mean(total_active_hours, na.rm = TRUE),
    avg_calories = mean(calories, na.rm = TRUE),
    avg_sleep_hours = mean(total_hours_asleep, na.rm = TRUE),
    sleep_efficiency = mean(total_minutes_asleep / total_time_in_bed, na.rm = TRUE)
  )
kpis
```


```
## # A tibble: 1 x 6
##   avg_steps median_steps avg_active_hours avg_calories avg_sleep_hours
##   <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1    7297.    7012.        3.64      2264.        7.00
## # i 1 more variable: sleep_efficiency <dbl>
```

Notes. Typical day shows moderate steps, ~6–8 hours sleep; calories track steps plus basal burn.

Average Steps by Weekday

```
avg_steps_weekday <- daily_activity %>%
  group_by(weekday) %>%
  summarise(avg_steps = mean(total_steps, na.rm = TRUE)) %>%
  ungroup()

ggplot(avg_steps_weekday, aes(x = weekday, y = avg_steps, fill = weekday)) +
  geom_col(show.legend = FALSE) +
  labs(title = "Average Steps by Weekday", x = "Weekday", y = "Average Steps") +
  theme_minimal()
```




Notes. Steps are steady across the week; slight midweek lift suggests routine-driven activity.

Steps vs Calories (Outliers Removed)

```
steps_cal_cor <- cor(daily_activity$total_steps, daily_activity$calories, use = "complete.obs")
steps_cal_cor
```

```
## [1] 0.5976419
```

```
ggplot(daily_activity, aes(x = total_steps, y = calories)) +
  geom_point(alpha = 0.45) +
  geom_smooth(method = "lm", se = FALSE) +
  labs(title = "Steps vs Calories Burned", x = "Total Steps", y = "Calories Burned") +
  theme_minimal()
```



Notes. Clear positive relationship; removing non-wear anomalies tightens the trend.

Steps × Sleep Hours → Mean Calories (Heatmap)

```
activity_sleep2 <- activity_sleep %>%
  mutate(
    total_hours_asleep = total_minutes_asleep / 60,
```


```

    is_weekend = factor(if_else(weekday %in% c("Saturday", "Sunday"), "Weekend", "Weekday"),
                        levels = c("Weekday", "Weekend"))
  )

y_max <- max(activity_sleep2$total_hours_asleep, na.rm = TRUE)
y_lim <- ceiling(y_max * 2) / 2
x_max <- quantile(activity_sleep2$total_steps, 0.99, na.rm = TRUE)

ggplot(activity_sleep2, aes(x = total_steps, y = total_hours_asleep)) +
  stat_summary_2d(aes(z = calories), fun = mean, bins = 30) +
  scale_fill_viridis(name = "Mean Calories", option = "C") +
  coord_cartesian(xlim = c(0, x_max)) +
  scale_y_continuous(breaks = seq(0, y_lim, 0.5)) +
  labs(title = "Steps vs Sleep Hours (color = mean calories)", x = "Total Steps", y = "Total Hours Asleep") +
  theme_minimal()


```



Notes. High calories concentrate with moderate-high steps and ~6–9 hours of sleep; very short sleep correlates with lower calories at similar steps.

Steps × Sleep Hours by Weekend/Weekday (Heatmap)

```
ggplot(activity_sleep2, aes(x = total_steps, y = total_hours_asleep)) +
  stat_summary_2d(aes(z = calories), fun = mean, bins = 30) +
  scale_fill_viridis(name = "Mean Calories", option = "C") +
  coord_cartesian(xlim = c(0, x_max)) +
  scale_y_continuous(breaks = seq(0, y_lim, 0.5)) +
  facet_wrap(~ is_weekend) +
  labs(title = "Steps vs Sleep Hours by Weekend/Weekday", x = "Total Steps", y = "Total Hours Asleep") +
  theme_minimal()
```



Notes. Weekend sleep is more variable; weekday patterns are tighter around workday schedules.

Active Hours by Weekday and Intensity

```
intensity_long <- daily_activity %>%
  select(weekday, very_active_minutes, fairly_active_minutes, lightly_active_minutes) %>%
  pivot_longer(
    cols = c(very_active_minutes, fairly_active_minutes, lightly_active_minutes),
    names_to = "intensity",
    values_to = "minutes"
```


```

) %>%
mutate(hours = minutes / 60)

y_max <- max(intensity_long$hours, na.rm = TRUE)
y_lim <- ceiling(y_max * 2) / 2

ggplot(intensity_long, aes(x = weekday, y = hours, fill = intensity)) +
  geom_col() +
  scale_y_continuous(breaks = seq(0, y_lim, 0.5)) +
  labs(title = "Active Hours by Weekday and Intensity", x = "Weekday", y = "Hours") +
  theme_minimal()

```



Notes. Light activity dominates; scope for micro-bursts of higher intensity without increasing total time burden.

Weekend vs Weekday (Goal & Distribution)

```

daily_activity <- daily_activity %>%
  mutate(
    is_weekend = factor(if_else(weekday %in% c("Saturday", "Sunday"), "Weekend", "Weekday"),
      levels = c("Weekday", "Weekend")),

```




```

    hit_10k = total_steps >= 10000
  )

pct_10k <- daily_activity %>%
  group_by(is_weekend) %>%
  summarise(pct = mean(hit_10k, na.rm = TRUE))

ggplot(pct_10k, aes(x = is_weekend, y = pct, fill = is_weekend)) +
  geom_col(show.legend = FALSE) +
  scale_y_continuous(labels = scales::percent) +
  labs(title = "Share of Days with 10,000 Steps", x = "", y = "Percent of Days") +
  theme_minimal()

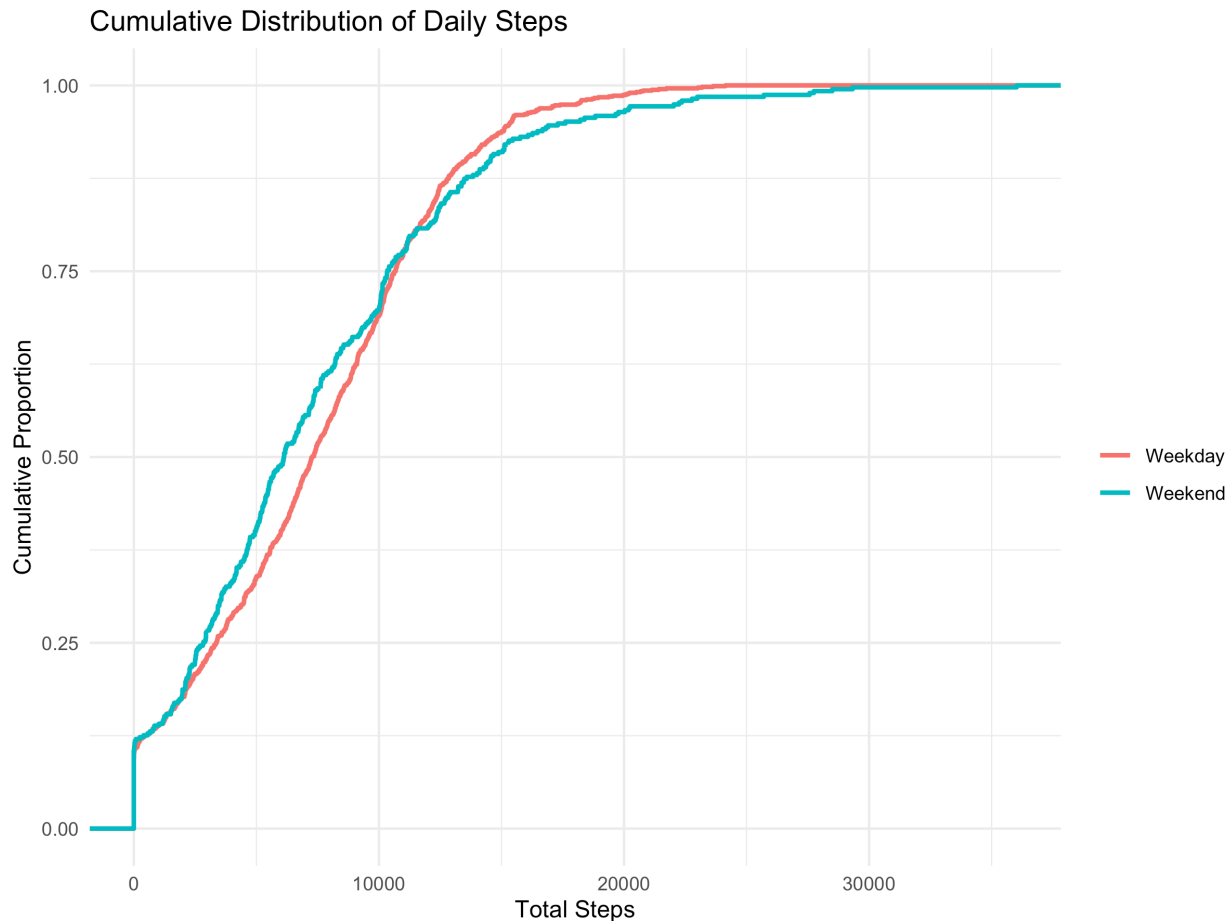
```



```

ggplot(daily_activity, aes(x = total_steps, color = is_weekend)) +
  stat_ecdf(size = 1) +
  labs(title = "Cumulative Distribution of Daily Steps", x = "Total Steps", y = "Cumulative Proportion") +
  theme_minimal()

```



Notes. Read proportions at 10k to quantify weekend vs weekday goal attainment.

Hourly Steps Heatmap

```
show_heatmap <- "hourlySteps_merged.csv" %in% names(data)

if (show_heatmap) {
  hourly_steps <- data[["hourlySteps_merged.csv"]] %>%
    janitor::clean_names() %>%
    mutate(
      activity_hour = mdy_hms(activity_hour),
      date = as.Date(activity_hour),
      hour = lubridate::hour(activity_hour),
      weekday = wday(date, label = TRUE, abbr = FALSE, week_start = 1)
    )

  heat <- hourly_steps %>%
    group_by(weekday, hour) %>%
    summarise(steps = mean(step_total, na.rm = TRUE), .groups = "drop")

  p08 <- ggplot(heat, aes(x = hour, y = weekday, fill = steps)) +
    geom_tile() +


```

```

labs(title = "Average Hourly Steps by Weekday", x = "Hour of Day", y = "Weekday", fill = "Avg Steps",
theme_minimal()

print(p08)
} else {
  message("hourlySteps_merged.csv not found in data; skipping heatmap.")
}

```



Notes - Step activity peaks during **morning (7–9 AM)** and **early evening (5–8 PM)** on most weekdays.
 - Weekends show a more diffuse pattern with weaker morning peaks. - Timing nudges around these windows should maximize engagement.

Share (Slide Deck Outline)

1. **Title Slide** — *Bellabeat Case Study – Fitbit Data Analysis*
 Tim Foltz – August 2025
2. **Business Task**
 - Analyze Fitbit usage trends to understand consumer habits

- Apply insights to Bellabeat products to improve marketing strategies
- Deliver actionable recommendations based on data findings

3. Data Sources

- Fitbit Fitness Tracker Dataset (CC0 Public Domain)
- ~30 participants with activity, sleep, and steps data
- Two periods: March–May 2016
- Some gaps in sleep and heart rate tracking

4. Methodology

- Data cleaning, standardization, and merging across date ranges
- Derived helper columns (weekend/weekday, total active hours)
- Filtered outliers (e.g., calories > 2500 with 0 steps)
- Explored patterns in activity, sleep, and calories

5. Plot 1 – Average Steps by Weekday

- Most steps midweek; Monday/Sunday lowest
- Suggests potential for weekend activity campaigns

6. Plot 2 – Steps vs Calories Burned

- Clear positive relationship
- Calories > 0 at ~0 steps suggests logged workouts without Fitbit movement data

7. Plot 3 – Sleep Hours vs Steps (Color/Size by Calories)

- High calories with moderate–high steps and ~6–9 hours sleep
- Very short sleep correlates with lower calories at similar steps

8. Plot 4 – Weekend vs Weekday % 10k Steps

- Weekdays higher very-active minutes; weekends more lightly active minutes

9. Plot 5 – Active Hours Intensity

- Light activity dominates; scope for short bursts of higher intensity without extra time

10. Plot 6 – Percent of Days with 10k Steps

- ~30% of days meet goal; no major weekday/weekend difference
- Longer weekend sleep may reduce active hours

11. Plot 7 – Distribution of Daily Steps

- ECDF shows similar weekend/weekday goal attainment proportions

12. Plot 8 – Hourly Steps Heatmap

- Moderate steps, 6–8 hours sleep, calories aligned with activity and basal burn
- Sleep efficiency stable; opportunity to raise fairly-active time

13. Key Findings

- Weekends lower average steps
- Steps strongly correlate with calories
- More sleep modestly linked to higher next-day activity
- Activity peaks midday with evening bump

14. Recommendations

- Weekend step challenges
- Sleep coaching for <7h sleepers
- Mixed-intensity badge rewards
- Hydration/activity nudges at peak hours

15. Act – Next Steps

- KPIs: +10% weekend steps, +5% active hours
- A/B test push notification timing
- Segment by sleep/activity patterns for tailored messaging
- Gather seasonal, demographic, and campaign-overlap data

Act (Recommendations & Next Steps)

Recommendations. - **10k Steps Challenge (all week).** Badges for streaks; spotlight top movers.

- **Sleep-Aware Coaching.** When sleep < 7h, suggest lighter goals and timed walk reminders.

- **Intensity Mix.** Promote 2×10 min fairly-active + 40 min lightly-active as a daily target.

Success Metrics. - Challenge participation rate; % days 10k; average steps delta vs baseline.

- Days with sleep <7h that still meet a scaled goal.

- Increase in fairly-active minutes without reducing total hours.

Experiment Plan. - A/B push timing (morning vs afternoon nudges).

- Targeting by prior-week sleep pattern.

- Content variants (badge vs streak messaging).

Data to Collect Next. - Longer time horizon (seasonality).

- Demographics and segments.

- Device wear-time to flag non-wear days.
- Campaign impression & click logs to tie to behavior change.

Limitations

Convenience sample (2016), potential non-wear and logging gaps, short window limits causal inference.

Reproducibility

```
set.seed(42)
sessionInfo()
```

```
## R version 4.5.1 (2025-06-13)
## Platform: aarch64-apple-darwin20
## Running under: macOS Tahoe 26.0
##
## Matrix products: default
## BLAS:   /Library/Frameworks/R.framework/Versions/4.5-arm64/Resources/lib/libRblas.0.dylib
## LAPACK: /Library/Frameworks/R.framework/Versions/4.5-arm64/Resources/lib/libRlapack.dylib; LAPACK v
##
## locale:
## [1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8
##
## time zone: America/Phoenix
## tzcode source: internal
##
## attached base packages:
## [1] stats      graphics  grDevices  utils      datasets  methods   base
##
## other attached packages:
## [1] viridis_0.6.5      viridisLite_0.4.2 scales_1.4.0      janitor_2.2.1
## [5] fs_1.6.6           lubridate_1.9.4   forcats_1.0.0    stringr_1.5.1
## [9] dplyr_1.1.4        purrr_1.0.4       readr_2.1.5      tidyr_1.3.1
## [13] tibble_3.3.0       tidyverse_2.0.0   ggplot2_3.5.2
##
## loaded via a namespace (and not attached):
## [1] generics_0.1.4     lattice_0.22-7    stringi_1.8.7     hms_1.1.3
## [5] digest_0.6.37      magrittr_2.0.3    evaluate_1.0.4    grid_4.5.1
## [9] timechange_0.3.0   RColorBrewer_1.1-3 fastmap_1.2.0     Matrix_1.7-3
## [13] gridExtra_2.3      mgcv_1.9-3        cli_3.6.5         rlang_1.1.6
## [17] crayon_1.5.3       splines_4.5.1     bit64_4.6.0-1     withr_3.0.2
## [21] yaml_2.3.10        tools_4.5.1       parallel_4.5.1    tzdb_0.5.0
## [25] vctrs_0.6.5        R6_2.6.1          lifecycle_1.0.4   snakecase_0.11.1
## [29] bit_4.6.0          vroom_1.6.5       pkgconfig_2.0.3   pillar_1.10.2
## [33] gtable_0.3.6       glue_1.8.0        xfun_0.52         tidyrselect_1.2.1
```

```
## [37] rstudioapi_0.17.1  knitr_1.50          farver_2.1.2         nlme_3.1-168
## [41] htmltools_0.5.8.1  rmarkdown_2.29       labeling_0.4.3       compiler_4.5.1
```